

---

## Chapter 6: Bayesian Learning

CS 536: Machine Learning  
Littman (Wu, TA)

---

### Roles for Bayesian Methods

Provides practical learning algorithms:

- Naive Bayes learning
- Bayesian belief network learning
- Combine prior knowledge (prior probabilities) with observed data
- Requires prior probabilities

Provides useful conceptual framework

- Provides “gold standard” for evaluating other learning algorithms
- Additional insight into Occam's razor

---

## Bayesian Learning

[Read Ch. 6, except 6.3]

[Suggested exercises: 6.1, 6.2, 6.6]

- Bayes Theorem
- MAP, ML hypotheses
- MAP learners
- Minimum description length principle
- Bayes optimal classifier
- Naive Bayes learner (if time)
- Example: Learning over text data
- Bayesian belief networks
- Expectation Maximization algorithm

---

### Bayes Theorem

$$P(h|D) = P(D|h) P(h) / P(D)$$

- $P(h)$  = prior prob. of hypothesis  $h$
- $P(D)$  = prior prob. of training data  $D$
- $P(h|D)$  = probability of  $h$  given  $D$
- $P(D|h)$  = probability of  $D$  given  $h$

## Choosing Hypotheses

---

Natural choice is most probable hypothesis given the training data, or *maximum a posteriori* hypothesis  $h_{MAP}$ :

$$\begin{aligned} h_{MAP} &= \operatorname{argmax}_{h \in H} P(h|D) \\ &= \operatorname{argmax}_{h \in H} P(D|h) P(h) / P(D) \\ &= \operatorname{argmax}_{h \in H} P(D|h) P(h) \end{aligned}$$

If assume  $P(h_i) = P(h_j)$  then can further simplify, and choose the *maximum likelihood* (ML) hypothesis

$$h_{ML} = \operatorname{argmax}_{h_i \in H} P(D|h_i)$$

## Basic Formulas for Probs

---

- Product Rule: probability  $P(A \wedge B)$  of a conjunction of two events A and B:  
 $P(A \wedge B) = P(A|B) P(B) = P(B|A) P(A)$
- Sum Rule: probability of a disjunction of two events A and B:  
 $P(A \vee B) = P(A) + P(B) - P(A \wedge B)$
- Theorem of total probability: if the events  $A_1, \dots, A_n$  are mutually exclusive with  $\sum_{i=1}^n P(A_i) = 1$ , then  
 $P(B) = \sum_{i=1}^n P(B|A_i) P(A_i)$

## Bayes Theorem

---

Does patient have cancer or not?

- A patient takes a lab test and the result comes back positive. The test returns a correct positive result in 98% of the cases in which the disease is actually present, and a correct negative result in 97% of the cases in which the disease is not present. Furthermore, .008 of the entire population have this cancer.

$$\begin{aligned} P(\text{cancer}) &= & P(\text{not cancer}) &= \\ P(+|\text{cancer}) &= & P(-|\text{cancer}) &= \\ P(+|\text{not cancer}) &= & P(-|\text{not cancer}) &= \end{aligned}$$


## Brute Force MAP Learner

---

1. For each hypothesis  $h$  in  $H$ , calculate the posterior probability  
$$P(h|D) = P(D|h) P(h) / P(D)$$
2. Output the hypothesis  $h_{MAP}$  with the highest posterior probability

$$h_{MAP} = \operatorname{argmax}_{h \in H} P(h|D)$$

## Evolution of Posterior Probs



- As data is added, certainty of hypotheses increases.
- What is the effect on entropy?

## Real-Valued Functions


Consider any real-valued target function  $f$   
 Training examples  $\langle x_i, d_i \rangle$ , where  $d_i$  is noisy training value

- $d_i = f(x_i) + e_i$
- $e_i$  is random variable (noise) drawn independently for each  $x_i$  according to some Gaussian distribution with mean=0

Then, the maximum likelihood hypothesis  $h_{ML}$  is the one that minimizes the sum of squared errors:

$$h_{ML} = \operatorname{argmin}_{h \in H} \sum_{i=1}^m (d_i - h(x_i))^2$$

## MAP and Least Squares



## MAP/Least Squares Proof

$$\begin{aligned} h_{MAP} &= \operatorname{argmax}_{h \in H} P(h|D) \\ &= \operatorname{argmax}_{h \in H} P(D|h) \\ &= \operatorname{argmax}_{h \in H} \prod_{i=1}^m \frac{1}{\sqrt{2\pi\sigma^2}} \cdot \exp\left(-\frac{1}{2} \left(\frac{d_i - h(x_i)}{\sigma}\right)^2\right) \\ &= \operatorname{argmax}_{h \in H} \sum_{i=1}^m \ln \frac{1}{\sqrt{2\pi\sigma^2}} - \frac{1}{2} \left(\frac{d_i - h(x_i)}{\sigma}\right)^2 \\ &= \operatorname{argmax}_{h \in H} \sum_{i=1}^m -\frac{1}{2} \left(\frac{d_i - h(x_i)}{\sigma}\right)^2 \\ &= \operatorname{argmax}_{h \in H} \sum_{i=1}^m -(d_i - h(x_i))^2 \\ &= \operatorname{argmin}_{h \in H} \sum_{i=1}^m (d_i - h(x_i))^2 \end{aligned}$$

## Predicting Probabilities

---

Consider predicting survival probability from patient data

Training examples  $\langle x_i, d_i \rangle$ , where  $d_i$  is 1 or 0

Want to train neural network to output a *probability* given  $x_i$  (not a 0 or 1)

## Predicting Probabilities

---

In this case, can show

$$h_{ML} = \operatorname{argmax}_{h \in H} \sum_{i=1}^m (d_i \ln h(x_i) + (1-d_i) \ln(1-h(x_i)))$$

Weight update rule for a sigmoid unit:

$$w_{jk} \leftarrow w_{jk} + \Delta w_{jk}$$

where  $\Delta w_{jk} = \eta \sum_{i=1}^m (d_i - h(x_i)) x_{ijk}$

## MDL Principle

---

Minimum Description Length Principle

Occam's razor: prefer the “shortest” hypothesis

MDL: prefer the hypothesis  $h$  that minimizes

$$h_{MDL} = \operatorname{argmin}_{h \in H} (L_{C1}(h) + L_{C2}(D|h))$$

where  $L_C(x)$  is the description length of  $x$  under encoding  $C$

## MDL Example

---

Example:  $H$  = decision trees,  $D$  = training data labels

- $L_{C1}(h)$  is # bits to describe tree  $h$
- $L_{C2}(D|h)$  is # bits to describe  $D$  given  $h$ 
  - Note  $L_{C2}(D|h) = 0$  if examples classified perfectly by  $h$ . Need only describe exceptions.
- Hence,  $h_{MDL}$  trades off tree size for training errors

## MDL Justification

---

$$\begin{aligned}h_{MAP} &= \operatorname{argmax}_{h \in H} P(D|h) P(h) \\&= \operatorname{argmax}_{h \in H} (\log_2 P(D|h) + \log_2 P(h)) \\&= \operatorname{argmin}_{h \in H} (-\log_2 P(D|h) - \log_2 P(h))\end{aligned}$$

From information theory:

The optimal (shortest expected coding length) code for an event with probability  $p$  is  $-\log_2 p$

So, prefer the hypothesis that minimizes  
 $length(h) + length(misclassifications)$

## Classification Example

---

Consider:

- Three possible hypotheses:  
 $P(h_1|D) = .4$ ,  $P(h_2|D) = .3$ ,  $P(h_3|D) = .3$
- Given new instance  $x$ ,  
 $h_1(x) = +$ ,  $h_2(x) = -$ ,  $h_3(x) = -$
- What's  $h_{MAP}(x)$  ?
- What's most probable classification of  $x$ ?

## Classifying New Instances

---

So far we've sought the most probable *hypothesis* given the data  $D$  (i.e.,  $h_{MAP}$ )

Given new instance  $x$ , what is its most probable *classification*?

- $h_{MAP}(x)$  is not the most probable classification!

## Bayes Optimal Classifier

---

**Bayes optimal classification:**

$$\operatorname{argmax}_{v_j \in V} \sum_{h_i \in H} P(v_j|h_i) P(h_i|D)$$

Example:

- $P(h_1|D) = .4$ ,  $P(-|h_1) = 0$ ,  $P(+|h_2) = 1$
  - $P(h_2|D) = .3$ ,  $P(-|h_2) = 1$ ,  $P(+|h_3) = 0$
  - $P(h_3|D) = .3$ ,  $P(-|h_3) = 1$ ,  $P(+|h_3) = 0$ ,
- therefore

$$\begin{aligned}\sum_{h_i \in H} P(+|h_i) P(h_i|D) &= .4 \\ \sum_{h_i \in H} P(-|h_i) P(h_i|D) &= .6 \quad \text{MAP class}\end{aligned}$$

## Gibbs Classifier

---

Bayes optimal classifier provides best result, but can be expensive if many hypotheses.

Gibbs algorithm:

1. Choose one hypothesis at random, according to  $P(h|D)$
2. Use this one to classify new instance

## Naive Bayes Classifier

---

Along with decision trees, neural networks, kNN, one of the most practical and most used learning methods.

When to use:

- Moderate or large training set available
- Attributes that describe instances are conditionally independent given classification

Successful applications:

- Diagnosis
- Classifying text documents

## Error of Gibbs

---

(Not so) surprising fact: Assume target concepts are drawn at random from  $H$  according to priors on  $H$ . Then:

$$E[\text{error}_{\text{Gibbs}}] \leq 2E[\text{error}_{\text{BayesOptimal}}]$$

Suppose correct, uniform prior distribution over  $H$ , then

- Pick any hypothesis consistent with the data, with uniform probability
- Its expected error no worse than twice Bayes optimal

## Naive Bayes Classifier

---

Assume target function  $f: X \rightarrow V$ , where each instance  $x$  described by attributes  $\langle a_1, a_2 \dots a_n \rangle$ .

Most probable value of  $f(x)$  is:

$$\begin{aligned} v_{MAP} &= \operatorname{argmax}_{v_j \in V} P(v_j | a_1, a_2 \dots a_n) \\ &= \operatorname{argmax}_{v_j \in V} \frac{P(a_1, a_2 \dots a_n, | v_j) P(v_j)}{P(a_1, a_2 \dots a_n)} \\ &= \operatorname{argmax}_{v_j \in V} P(a_1, a_2 \dots a_n, | v_j) P(v_j) \end{aligned}$$

## Naïve Bayes Assumption

---

$$P(a_1, a_2 \dots a_n | v_j) = \prod_i P(a_i | v_j),$$

which gives

**Naïve Bayes classifier:**

$$v_{NB} = \operatorname{argmax}_{v_j \text{ in } V} P(v_j) \prod_i P(a_i | v_j)$$

Note: No search in training!

## Naïve Bayes Algorithm

---

Naïve\_Bayes\_Learn(*examples*)

For each target value  $v_j$

$$\hat{P}(v_j) \leftarrow \text{estimate } P(v_j)$$

For each attribute value  $a_i$  of each attribute  $a$

$$\hat{P}(a_i | v_j) \leftarrow \text{estimate } P(a_i | v_j)$$

Classify\_New\_Instance( $x$ )

$$v_{NB} = \operatorname{argmax}_{v_j \text{ in } V} \hat{P}(v_j) \prod_i \hat{P}(a_i | v_j)$$

## Naïve Bayes: Example

---

- Consider *PlayTennis* again, and new instance

*<Outlk = sun, Temp = cool, Humid = high, Wind = strong>*

Want to compute:

$$v_{NB} = \operatorname{argmax}_{v_j \text{ in } V} P(v_j) \prod_i P(a_i | v_j)$$

$$P(y) P(\text{sun}|y) P(\text{cool}|y) P(\text{high}|y) P(\text{strong}|y) = .005$$

$$P(n) P(\text{sun}|n) P(\text{cool}|n) P(\text{high}|n) P(\text{strong}|n) = .021$$

- So,  $v_{NB} = n$

## Naïve Bayes: Subtleties

---

1. Conditional independence assumption is often violated

$$P(a_1, a_2 \dots a_n | v_j) = \prod_i P(a_i | v_j)$$

- ...but it works surprisingly well anyway. Note don't need estimated posteriors  $P(v_j|x)$  to be correct; need only that

$$\begin{aligned} \operatorname{argmax}_{v_j \text{ in } V} P(v_j | a_1, a_2 \dots a_n) \\ = \operatorname{argmax}_{v_j \text{ in } V} P(v_j) \prod_i P(a_i | v_j) \end{aligned}$$

- Domingos & Pazzani [1996] for analysis
- Naïve Bayes posteriors often unrealistically close to 1 or 0

## Naïve Bayes: Subtleties

---

2. what if none of the training instances with target value  $v_j$  have attribute  $a_i$ ?

$$P(a_i|v_j) = 0, \text{ and... } P(v_j) \prod_i P(a_i|v_j) = 0$$

Solution is Bayesian estimate:

$$P(a_i|v_j) = (n_c + mp)/(n + m) \text{ where}$$

- $n$  is number of training examples for which  $v = v_j$ ,
- $n_c$  number of examples for which  $v = v_j$  and  $a = a_i$
- $p$  is prior estimate for  $P(a_i|v_j)$
- $m$  is weight given to prior (i.e., number of "virtual" examples)